

PDE-Agents: An LLM-Orchestrated Multi-Agent Framework for Automated Finite Element Simulations with Knowledge Graph-Augmented Reasoning

Sayan Adhikari, Gulshan Noorsumar, Øyvind Jensen

MatPro, Institute for Energy Technology

[email]

Submitted: June 5, 2026

Abstract

We present *PDE-Agents*, a multi-agent ecosystem that automates the full lifecycle of partial differential equation (PDE) / finite element method (FEM) simulations through natural-language interaction. The system combines three specialist large language model (LLM) agents—Simulation, Analytics, and Database—orchestrated via a LangGraph supervisor graph, with a locally-deployed open-source LLM stack (Qwen3-Coder-Next, Llama 4 Scout) running on dual NVIDIA RTX PRO 6000 Blackwell GPUs (≈ 196 GB combined VRAM). The architecture is model-agnostic; we report cross-model validation across two generations of open-source LLMs. A central novelty is the integration of a *GraphRAG* knowledge base (Neo4j + 768-dimensional vector embeddings) that enriches agents with curated material properties, known failure patterns, and prior run lineage. We report seven empirical contributions: (i) a formal verification and validation (V&V) study confirming second-order spatial convergence ($\mathcal{O}(h^2)$) for all three benchmark cases on the heat equation solver; (ii) a three-way ablation study across 50 benchmark tasks with a frozen KG, comparing KG On, KG Off, and KG Smart, showing that KG Smart achieves 100% success *and* the highest output quality (physics score 0.933 vs. 0.853 for KG Off, MPF 0.926 vs. 0.796), and that our *KG Smart* integration, combining warm-start injection with lazy conditional retrieval, matches KG Off reliability while maximising output fidelity; (iii) a novel-material experiment using three fictional materials whose properties exist only in the KG, where KG Smart achieves near-perfect material property fidelity (MPF = 1.00) versus 0.34 for the KG-free baseline; (iv) a failure analysis tracing KG On’s 3 systematic failures to budget exhaustion and timeout, establishing warm-start injection as the dominant factor in KG Smart’s reliability advantage; (v) an adaptive KG decision framework (Algorithm 1) that selects the optimal retrieval mode per task; (vi) production-scale agent quality metrics from 1,369 real simulation runs showing 97.8% overall success and a 57.6% first-try rate; and (vii) a controlled 100-task KG growth experiment showing that accumulated run history yields a difficulty-dependent quality gain, with hard-task

MPF improving by 8.8% between passes while easy and novel tasks remain at ceiling. All code, models, and evaluation artifacts are released openly. Our findings demonstrate that *integration pattern*, not knowledge content, determines whether GraphRAG augmentation helps or hinders LLM agents, and identify concrete design principles for autonomous FEM simulation assistants.

Keywords: Large language model agents; finite element method; knowledge graph; GraphRAG; multi-agent systems; scientific computing automation; verification and validation.

1 Introduction

Finite element method simulations are central to engineering analysis across structural mechanics, heat transfer, fluid dynamics, and electromagnetics. Yet the gap between a domain expert’s intent—“what if I use aluminium instead of steel?”—and a running, verified simulation involves a sequence of error-prone manual steps: geometry creation, boundary condition specification, solver parameter selection, result interpretation, and iterative debugging.

The recent emergence of *reasoning-capable* LLMs [1–3] has opened the possibility of agents that can autonomously traverse this workflow. Prior work on “AI for scientific computing” has largely focused on surrogate modelling [4–6], operator learning [7], or dataset-specific fine-tuning [8]. Much less attention has been paid to *procedural automation*—using LLMs to *set up and manage* solvers rather than to replace them.

Meanwhile, the *retrieval-augmented generation* (RAG) paradigm [9] has demonstrated that grounding LLM outputs with authoritative external knowledge dramatically reduces hallucination. The graph-structured variant, GraphRAG [10], further enables multi-hop reasoning over relational knowledge—particularly relevant in engineering contexts where material properties, mesh requirements, and solver stability rules form a rich, interconnected knowledge base.

In this paper we make the following contributions:

1. We design and implement *PDE-Agents*, a complete,

containerised multi-agent ecosystem for automated FEM simulation (section 3).

2. We describe a novel *GraphRAG integration pattern* that encodes material properties, known failure modes, and prior run lineage as a Neo4j knowledge graph with Hierarchical Navigable Small World (HNSW) vector-indexed embeddings (section 4).
3. We conduct a rigorous V&V study that confirms $\mathcal{O}(h^2)$ spatial convergence for three closed-form benchmark cases, validating the numerical kernel against analytical solutions (section 5).
4. We run a three-way ablation study (KG On, KG Off, KG Smart) over 50 tasks with a frozen knowledge graph, demonstrating that KG-enabled modes produce higher-quality simulation outputs (physics score, material property fidelity) than the KG-free baseline, and that *KG Smart* achieves 100% success and the highest output quality (section 6).
5. We introduce a *novel-material experiment* using three fictional materials absent from any LLM training corpus, demonstrating that KG Smart achieves $\approx 2.9\times$ higher material property fidelity than the KG-free baseline, which fabricates physically incorrect properties (section 6.4).
6. We trace KG On’s 3 systematic failures to budget exhaustion and timeout, revealing that warm-start injection, not lazy retrieval, is the dominant factor in KG Smart’s reliability advantage (section 6.5).
7. We report production-scale agent quality metrics derived from 1,369 real simulation runs and a controlled 100-task KG growth experiment demonstrating a difficulty-dependent quality gain: hard-task MPF improves by 8.8% as the KG accumulates run history (section 7).

2 Related Work

LLM agents for scientific computing. ChemCrow [11] demonstrated that equipping an LLM with chemistry tools (RDKit, reaction planners) enables autonomous laboratory-scale reasoning. SciAgent [12] and similar systems extend this to broader scientific question answering. Our work targets *numerical simulation* rather than knowledge retrieval, requiring tool-call pipelines that produce and verify deterministic numerical outputs.

Autonomous FEM systems. FEniCS [13] and its successor DOLFINx/FEniCSx [14] provide Python-scriptable FEM solvers that are amenable to LLM-driven automation. Early work on “simulation copilots” [15, 16] shows that LLMs can generate solver scripts from natural language with moderate success, but these approaches are single-turn and lack closed-loop debugging. More recently, ALL-FEM [17] fine-tunes LLMs on 1,000+ verified

FEniCS scripts and embeds them in a multi-agent workflow achieving 71.8% code-level success on multiphysics benchmarks, and FEABench [18] provides a systematic evaluation using COMSOL Multiphysics (88% executable API calls). PDE-Agents differs by using unmodified open-source LLMs with knowledge-graph augmentation rather than domain-specific fine-tuning, and by providing a controlled ablation of retrieval strategies with physics-aware quality metrics.

Multi-agent orchestration. LangGraph [19] implements graph-structured agent workflows that support conditional routing and persistent state, enabling the supervisor pattern we adopt. AutoGen [20] and CrewAI [21] offer related frameworks; we choose LangGraph for its explicit state graph semantics and tight LangChain integration.

Retrieval-augmented generation for engineering. RAG [9] has been applied to engineering documentation retrieval [22], code generation [23], and more recently to simulation parameter suggestion [24]. GraphRAG [10] adds structured relational traversal. Corrective RAG (CRAG) [25] introduces adaptive retrieval decisions based on document relevance scoring, and AriGraph [26] demonstrates episodic knowledge-graph memory for LLM agents. Our *KG Smart* pattern draws on both: warm-start injection from similar past runs (episodic memory) with lazy conditional tool invocation (corrective retrieval). To our knowledge, ours is the first system to combine GraphRAG with a live, iterative FEM simulation loop.

Knowledge graphs for materials science. Materials-focused knowledge graphs (e.g. MatKG [27], OPTIMADE [28]) encode material properties in RDF/OWL formats. Buehler [29] demonstrates that retrieval-augmented ontological KG strategies outperform flat RAG for materials design by capturing mechanistic relationships between concepts. Our Neo4j KG is lighter-weight and purpose-built for run-time agent queries, prioritising latency over completeness.

3 System Architecture

3.1 Design Philosophy

PDE-Agents adopts four guiding principles:

1. **Open and local:** All models run locally via Ollama on dual NVIDIA RTX PRO 6000 Blackwell GPUs (≈ 196 GB VRAM, CUDA 13.1), eliminating API costs and data-privacy concerns.
2. **Verifiable:** Every agent reasoning step, tool call, and result is logged to a relational database, enabling post-hoc audit.
3. **Containerised:** All 11 services are orchestrated via Docker Compose, ensuring reproducible deployment.




Figure 1: System architecture of PDE-Agents (four-tier layout). *Tier 1*: user natural-language interface. *Tier 2*: Multi-Agent Orchestration — a LangGraph Supervisor routes tasks to three specialist agents (Simulation, Analytics, Database), each backed by a locally-deployed LLM via Ollama. *Tier 3*: Execution & Knowledge — FEniCSx FEM runner (DOLFINx 0.10.0.post2), Neo4j GraphRAG Knowledge Graph (HNSW-indexed 768-d embeddings), and persistent storage (PostgreSQL + MinIO). *Tier 4*: shared Ollama LLM Backend (196 GB VRAM, CUDA 13.1). The thick dashed red arrow (Simulation Agent $\rightarrow$ Neo4j) represents the *KG Smart warm-start*: the top-3 similar past runs are retrieved via HNSW vector search and injected into the agent’s prompt before the reasoning loop. The thin dashed red arrow (Analytics Agent $\rightarrow$ Neo4j) is the standard lazy KG query path. The blue bidirectional arrow at right carries LLM inference traffic between the orchestration layer and the shared Ollama backend.

4. **Separation of concerns:** The FEM solver, LLM stack, knowledge graph, and persistence layer are decoupled services communicating over a private bridge network.

3.2 Agent Stack

Each agent’s LLM backend is configurable via environment variables, enabling drop-in model upgrades without code changes. table 1 lists the default models and the alternatives validated in the cross-model experiments of section 6.4.

Table 1: Agent model configuration. Current defaults reflect the model upgrade from the initial (v1) deployment. The ablation study uses a frozen KG snapshot across 50 benchmark tasks per mode.

Agent	Env Variable	Default (current)	Previous (v1)
Orch.	ORCHESTRATOR_MODEL	llama4:scout	llama3.3:70b
Sim.	SIM._AGENT_MODEL	qwen3-coder-next	qwen2.5-coder:32b
Anal.	ANAL._AGENT_MODEL	llama4:scout	llama3.3:70b
DB	DB_AGENT_MODEL	qwen3-coder:30b	qwen2.5-coder:14b

Orchestrator. A LangGraph StateGraph supervisor receives the user’s natural-language request and routes it to one of three specialist agents via structured JSON decisions. After each agent reports its result the supervisor either routes to another agent or synthesises a final response.

Simulation Agent. The core reasoning agent, executing a ReAct loop [3] with up to 25 reasoning steps and nine tools: `check_config_warnings`, `query_knowledge_graph`, `validate_config`, `run_simulation`, `modify_config`, `debug_simulation`, `list_recent_runs`, `get_run_status`, `run_parametric_sweep`.

Analytics Agent. Max 12 iterations; tools for statistical comparison across runs, sensitivity analysis, and LLM-generated textual summaries.

Database Agent. Max 10 iterations; answers history queries against PostgreSQL, cataloguing results and tracing run lineage.

3.3 Tool-Call Compatibility

Supporting multiple LLM families requires robust tool-call parsing, as models encode tool invocations differently. The Qwen 2.5-Coder family emits tool calls as JSON text in the `content` field rather than in the structured `tool_calls` field expected by LangChain’s ToolNode. We implement a four-pass normalisation procedure in `BaseAgent._parse_content_tool_call`: (1) attempt `json.loads` on the full content; (2) extract from markdown fences; (3) extract the first `{...}` block; (4) repair truncated JSON via progressive bracket-closing and regex fallback.

A separate compatibility issue emerged with Qwen3-Coder-Next: the model occasionally *double-encodes* JSON tool arguments (e.g. `{"key": "val"}`). All tool functions now use a `_safe_json_parse` helper that detects and recursively unwraps such double-encoded strings, ensuring seamless operation across model families without model-specific workarounds. Llama 4 Scout, by contrast, uses the standard `tool_calls` field natively and requires no special parsing.

3.4 FEM Solver: DOLFINx Heat Equation

The FEniCSx runner container (DOLFINx 0.10.0.post2) exposes a FastAPI endpoint that accepts a `HeatConfig` JSON and returns a result including run metadata. The solver uses P1 (linear Lagrange) elements, with Backward Euler ($\theta = 1$, unconditionally stable) as the default time integration scheme, supporting Dirichlet, Neumann, and Robin boundary conditions. Geometry is built with Gmsh [30] for 2D and 3D domains from a library of nine parameterised types: `rectangle`, `l_shape`, `circle`, `annulus`, `hollow_rectangle`, `t_shape`, `stepped_notch`, `box`, and `cylinder`; Gmsh geometries use named physical groups for boundary identification, enabling arbitrary complex domain topologies.

4 Knowledge Graph Architecture

The Neo4j knowledge graph (Neo4j 5 Community Edition) serves three functions:

1. **Material property lookup.** A curated library of materials (metals, ceramics, polymers) with thermal conductivity k , density ρ , and specific heat c_p ranges, sourced from engineering handbooks and validated against ASM International data. Approximate values: steel $k \approx 50$ W/mK, copper $k \approx 385$ W/mK, aluminium $k \approx 200$ W/mK.
2. **Failure pattern catalogue.** `KnownIssue` nodes encode failure modes observed during system operation. A pure-Python rule engine (`rules.py`) fires nine pre-run checks analytically before each simulation:

Rule code	Trigger
INCONSISTENT_IC	$ u_{\text{init}} - T_{\text{BC,min}} > 100 \text{ K}$
EXPLICIT_CFL_VIOLATION	$\theta < 0.5$ and $\Delta t > h^2/(2\alpha)$
NEAR_EXPLICIT_SCHEME	$0 \leq \theta < 0.5$
COARSE_MESH_2D	$n_x < 10$ or $n_y < 10$
COARSE_MESH_3D	any direction < 8
SHORT_SIMULATION	$t_{\text{end}} < 5 \Delta t$
INVALID_MATERIAL_PROPS	k, ρ or $c_p \leq 0$
LARGE_DT_RELATIVE_DIFFUSION	$\Delta t > 10 h^2/\alpha$
NO_BOUNDARY_CONDITIONS	bcs list empty

3. **Run lineage and semantic retrieval.** Each completed `Run` node is embedded with `nomic-embed-text` (768 dimensions) and indexed




Figure 2: Knowledge graph visualisation (Neo4j-style). **Blue circles:** material nodes (double border = novel/fictional materials known only to the KG). **Green circles:** simulation run nodes linked to their materials via **USES_MATERIAL** edges and to each other via dashed **SIMILAR_To** edges (cosine similarity on 768-d HNSW embeddings; scores shown on edges). **Orange diamonds:** known-issue nodes triggered by failed runs. **Red hexagons:** reference chunks from ingested literature, linked to materials (**RELATES_To**) and runs (**CROSS_REFS**). **Grey boxes:** boundary-condition pattern nodes. Pill-shaped nodes show the thermal-class taxonomy.

in a Hierarchical Navigable Small World (HNSW) vector index. Up to $k=5$ `SIMILAR_TO` edges are created to the nearest neighbours with cosine similarity ≥ 0.85 , precomputing the neighbourhood graph for fast agent traversal. The Simulation Agent can query “find similar past runs” to retrieve configurations with known outcomes, enabling few-shot in-context learning from the system’s own history.

Document intelligence pipeline. A Celery-backed ingestion pipeline processes PDF papers, standards documents, and HTML references using Docling [31] for structured extraction, splitting them into `ReferenceChunk` nodes with context-aware 256-token windows (overlap 32 tokens) and HNSW-indexed embeddings. Each chunk is automatically cross-referenced to semantically similar `Run` nodes via `CROSS_REFS` edges (cosine ≥ 0.78), bridging literature knowledge to simulation history. As a representative scale point: indexing the 50-page FEniCSx tutorial produced 298 chunks and 1,073 cross-references.

5 Verification and Validation

We conduct a formal V&V study following ASME V&V 10-2006 guidelines [32], comparing the DOLFINx FEM solver against closed-form analytical solutions on three benchmark cases.

5.1 Benchmark Cases

Case 1: 2D Steady-State Linear Profile. Laplace equation on $\Omega = [0, 1]^2$ with Dirichlet conditions $T|_{x=0} = 0$, $T|_{x=1} = 1$ and Neumann zero-flux on top/bottom. Analytical solution: $T(x, y) = x$. P1 elements are nodally exact for linear solutions; we use UFL spatial-coordinate expressions in high-order quadrature (degree 8) to measure the true continuous L^2 error.

Case 2: 2D Transient Fourier Mode Decay. Heat equation $\partial_t T = \alpha \nabla^2 T$ with homogeneous Dirichlet on all boundaries and initial condition $T_0 = \sin(\pi x) \sin(\pi y)$. Analytical solution: $T(x, y, t) = e^{-2\alpha\pi^2 t} \sin(\pi x) \sin(\pi y)$.

Case 3: 2D Steady-State Poisson with Constant Source. $-k \nabla^2 T = f$ with $f = 1000 \text{ W/m}^3$, Dirichlet $T = 0$ on left/right, Neumann on top/bottom. Exact solution: $T(x) = f(x - x^2)/(2k)$.

5.2 Convergence Study Protocol

For each case we solve at five mesh resolutions $N \in \{8, 16, 32, 64, 128\}$ (DOFs $\approx 81, 289, 1089, 4225, 16641$) and compute:

$$\|e\|_{L^2} = \left(\int_{\Omega} (u_h - u_{\text{exact}})^2 dx \right)^{1/2} \quad (1)$$

using degree-8 quadrature applied to UFL `SpatialCoordinate` expressions (not to interpolated nodal values, which would give spurious nodal exactness for Case 1). The convergence rate is estimated by linear regression on the log-log plot.




Figure 3: Spatial convergence study. Cases 2 and 3 exhibit the expected $\mathcal{O}(h^2)$ rate for P1 elements. Case 1 is algebraically exact (linear profile is represented exactly in the P1 space), with $\|e\|_{L^2} \approx \mathcal{O}(10^{-9})$ attributable only to floating-point rounding.

5.3 Results

All three cases pass: Cases 2 and 3 achieve rates of 2.04 and 2.00, consistent with the theoretical $\mathcal{O}(h^2)$ prediction for P1 finite elements on uniform Cartesian meshes. Case 1 achieves machine-precision errors (order 10^{-9}) because the linear profile lies exactly in the P1 polynomial space, confirming correct solver implementation.

5.4 Representative Simulation Gallery

Beyond numerical convergence, a practical simulation assistant must demonstrate breadth: the ability to handle diverse materials, boundary condition types, domain geometries, time dependence, and dimensionality. Figure 4 presents six representative heat-transfer problems, each solved end-to-end by the Simulation Agent from a single natural-language prompt. No manual intervention was required for mesh generation, solver configuration, boundary condition application, or post-processing. All cases use P1 Lagrange elements on triangular (2D) or tetrahedral (3D) meshes, solved by DOLFINx 0.10.0.post2. Temperature fields are reported in Kelvin throughout for consistency.

Materials. The gallery covers four distinct engineering alloys whose thermal conductivities span two orders of magnitude: copper ($k = 385 \text{ W/(m}\cdot\text{K)}$, panel a), AISI 1010 steel ($k = 50 \text{ W/(m}\cdot\text{K)}$, panel b), aluminium 6061 ($k = 205 \text{ W/(m}\cdot\text{K)}$, panels d and e), Ti-6Al-4V titanium ($k = 6.7 \text{ W/(m}\cdot\text{K)}$, panel c), and stainless steel 304 ($k = 16.3 \text{ W/(m}\cdot\text{K)}$, panel f). In a production deployment, the

Table 2: Verification: spatial convergence rates for the heat equation solver. Expected rate is $\mathcal{O}(h^2)$ for P_1 elements.

Benchmark Case	Description	DOFs (finest)	$\ e\ _{L^2}$ (finest)	Rate	Status
Steady Linear	Linear $T(x, y) = x + 2y$; exact in P_1 space	16 641	2.47×10^{-9}	<i>exact</i>	✓
Transient Fourier	$e^{-2\pi^2 t} \sin(\pi x) \sin(\pi y)$ decay mode	16 641	1.50×10^{-5}	2.04	✓
Steady Poisson	$-\nabla^2 u = 1$; exact soln. via Green’s fn.	16 641	5.57×10^{-6}	2.00	✓

agent retrieves these values from the Neo4j knowledge graph via `query_knowledge_graph`.

Boundary conditions. Panel (a) applies the simplest configuration: Dirichlet conditions on opposing faces ($T = 373.15$ K left, $T = 273.15$ K right), producing the expected linear gradient. Panel (b) combines all four major BC types on a single domain: Dirichlet (300 K, left), Neumann inward flux (2 kW/m^2 , top), Robin convection ($h = 25 \text{ W/(m}^2 \cdot \text{K)}$, $T_\infty = 293$ K, right), and insulated (bottom), together with volumetric heating $Q = 5 \text{ kW/m}^3$, creating the asymmetric field visible in the plot. Panel (c) uses Dirichlet BCs on opposing edges (573 K left, 293 K right) with the insulated hole boundary distorting the field into characteristic thermal concentration bands. Panel (f) combines Dirichlet (300 K bottom, 600 K left face) with Robin convection on the top face of a 3D cube.

Geometry. Panels (a), (b), and (d) use structured Cartesian meshes on the unit square. Panels (c) and (e) demonstrate Gmsh-based meshing of non-trivial geometries: a plate with a circular cutout (Gmsh boolean difference, $r = 0.15$, ≈ 5000 vertices) and an L-shaped domain ($[0, 1]^2 \setminus [0.5, 1]^2$, mesh size ≈ 0.02). Panel (f) meshes a unit cube with 24^3 tetrahedra (≈ 15600 vertices), demonstrating 3D capability.

Time dependence. Panel (e) solves the transient heat equation using 100 implicit-Euler steps ($\Delta t = 0.005$ s) on the L-shaped domain (aluminium, $\rho = 2700 \text{ kg/m}^3$, $c_p = 900 \text{ J/(kg} \cdot \text{K)}$), with the snapshot at $t = 0.50$ s showing the thermal front propagating from the hot left edge.

Source terms. Panel (d) features a localised Gaussian volumetric heat source $Q(x, y) = Q_{\text{peak}} \exp\left(-\frac{(x-x_0)^2 + (y-y_0)^2}{2\sigma^2}\right)$ with $Q_{\text{peak}} = 5 \times 10^6 \text{ W/m}^3$, centred at $(0.3, 0.7)$ with $\sigma = 0.08$, and Dirichlet boundaries at 293 K. This setup is representative of laser spot heating and Joule-heating applications, and produces the characteristic concentric isotherms visible in the figure.

Reproducibility. Each case is implemented as a self-contained Python script in `evaluation/examples/` (see Table 3). All simulation parameters (material properties, mesh resolution, boundary values, time stepping) are defined at the top of each script, making it straightforward for reviewers to modify and re-run individual cases.

Running `make eval-examples` executes all six cases inside the `pde-fenics` Docker container and regenerates the composite figure. Intermediate results are saved as NumPy `.npz` archives containing mesh coordinates, cell connectivity, and solution vectors, enabling figure regeneration without re-running the solver.

Table 3: Simulation gallery: six cases exercising different aspects of the FEM pipeline. Scripts in `evaluation/examples/`.

	Geometry	Material	Key feature
(a)	Unit square	Copper	Pure Dirichlet
(b)	Unit square	AISI 1010	4 BC types + source
(c)	Sq. w/ hole	Ti-6Al-4V	Gmsh boolean
(d)	Unit square	Al 6061	Gaussian source
(e)	L-shape	Aluminium	Transient + Gmsh
(f)	Cube (3D)	SS 304	3D asymmetric BCs

6 Ablation Study: Knowledge Graph Contribution

6.1 Experimental Design

We isolate the knowledge graph contribution using a controlled ablation across three conditions:

- **KG On** (mandatory): Full system with `check_config_warnings` and `query_knowledge_graph` tools enabled; the system prompt requires their use before every run.
- **KG Off** (baseline): Identical system with KG tools removed entirely; the agent proceeds directly to validation and execution.
- **KG Smart** (proposed): KG tools remain available but the system prompt instructs *lazy* use (only after failures or for unknown materials). Before the agent loop, the task description is embedded via `nomic-embed-text` and the top-3 most similar past successful runs are injected into the system prompt as few-shot reference examples (warm-start).

fig. 5 contrasts the agent workflow under each condition. KG Off proceeds directly from parsing to simulation; KG On forces two KG round-trips before every attempt (and loops on failure); KG Smart front-loads a single HNSW warm-start and defers KG queries to the failure-recovery path only.




Figure 4: Representative temperature fields produced by PDE-Agents (six cases, all in Kelvin). **(a)** Steady-state Dirichlet on copper ($k = 385 \text{ W}/(\text{m}\cdot\text{K})$; $T = 373.15 \text{ K}$ left, $T = 273.15 \text{ K}$ right; $N = 64$). **(b)** Mixed BCs on AISI 1010 steel ($k = 50 \text{ W}/(\text{m}\cdot\text{K})$): Dirichlet 300 K left, Neumann $2 \text{ kW}/\text{m}^2$ top, Robin $h = 25 \text{ W}/(\text{m}^2\cdot\text{K})$ right, insulated bottom, $Q = 5 \text{ kW}/\text{m}^3$. **(c)** Plate with circular hole ($r = 0.15$) in Ti-6Al-4V ($k = 6.7 \text{ W}/(\text{m}\cdot\text{K})$); Gmsh boolean difference; Dirichlet $573 \text{ K}/293 \text{ K}$ on opposing edges. **(d)** Localised Gaussian heat source on Al 6061 ($Q_{\text{peak}} = 5 \times 10^6 \text{ W}/\text{m}^3$, $\sigma = 0.08$, centre $(0.3, 0.7)$); Dirichlet 293 K on all boundaries. **(e)** Transient L-shaped domain in aluminium at $t = 0.50 \text{ s}$ (implicit Euler, $\Delta t = 0.005 \text{ s}$, 100 steps). **(f)** 3D unit cube in SS 304 ($k = 16.3 \text{ W}/(\text{m}\cdot\text{K})$), horizontal cross-section at $z = 0.5$; Dirichlet bottom 300 K and left face 600 K , Robin top. All solved with DOLFINx 0.10.0.post2, P1 elements. Reproducible scripts in `evaluation/examples/`.




Figure 5: Agent workflow under the three KG integration modes. **(a)** KG Off: fast but the LLM fabricates material properties. **(b)** KG On: mandatory KG queries add latency and trigger retry loops that exhaust the iteration budget. **(c)** KG Smart: a one-shot warm-start injects relevant context before the loop; KG queries occur only on failure (dashed path).

Methodology. To eliminate data leakage between conditions, we *freeze* the knowledge graph (set `KG_READ_ONLY=true` at the code level) before starting the ablation: all three modes read from the same KG snapshot, but no mode can write to it, ensuring that a failure in one mode cannot help subsequent modes. Tasks are shuffled independently per mode and run in-process inside the agents container to eliminate HTTP overhead. Confidence intervals use Wilson scores; pairwise comparisons use chi-squared tests and Cohen’s h / Cohen’s d effect sizes.

Benchmark tasks. Fifty tasks span four difficulty levels:

- **Easy (12 tasks):** Explicit numerical parameters, straightforward setup.
- **Medium (15 tasks):** Material-by-name requests requiring property lookup (e.g. steel, copper, aluminium, zinc).
- **Hard (13 tasks):** Ambiguous descriptions, mixed BCs, computationally expensive or 3D problems.
- **Novel (10 tasks):** Tasks referencing three fictional materials whose thermal properties exist *only* in the knowledge graph, designed to stress-test the agent’s ability to retrieve domain knowledge rather than rely on LLM memorisation.

The three fictional materials are:

- **Novidium:** a ceramic-metallic composite ($k=73$ W/m K, $\rho=5420$ kg/m³, $c_p=612$ J/kg K), a moderate conductor with unusual density.
- **Cryonite:** a polymer-aerogel hybrid ($k=0.42$, $\rho=1180$, $c_p=1940$), an extreme insulator with very low diffusivity ($\alpha \approx 1.8 \times 10^{-7}$ m²/s).
- **Pyrrathane:** a refractory cermet ($k=312$, $\rho=3850$, $c_p=278$), exhibiting very high conductivity and diffusivity ($\alpha \approx 2.9 \times 10^{-4}$ m²/s).

Each material is seeded as a **Material** node in Neo4j with linked **Reference** chunks containing exact property values and simulation guidance. The 10 novel tasks span steady-state, transient, and mixed-BC configurations across all three materials.

Evaluation metrics. Standard success rate is inadequate for evaluating KG utility because a simulation with fabricated properties still “succeeds” computationally. We therefore introduce two physics-aware metrics:

- **Material Property Fidelity (MPF):** $MPF = 1 - \frac{1}{3} \sum_{p \in \{k, \rho, c_p\}} |p_{\text{agent}} - p_{\text{truth}}| / p_{\text{truth}}$, measuring how accurately the agent retrieves the correct material constants.

- **Physics Score:** $0.5 \cdot \text{MPF} + 0.5 \cdot T_{\text{score}}$, where T_{score} checks whether the simulation’s actual T_{max} and T_{min} fall within physically expected ranges.
- **Sensitivity-weighted MPF (MPF_w):** $\text{MPF}_w = 1 - \sum_p S_p \cdot |p_{\text{agent}} - p_{\text{truth}}| / (p_{\text{truth}} \cdot \sum_q S_q)$, where $S_p = |\partial T_{\text{mean}} / \partial p|$ are sensitivity coefficients computed via central finite differences (section 6.4). For steady-state Dirichlet tasks where $S_p \approx 0$, we set $\text{MPF}_w=1$ since wrong properties have no output impact. This metric penalises errors in high-sensitivity properties (e.g. k in transient problems) more heavily than errors in low-sensitivity properties (e.g. ρ in steady-state).

On the 10 novel tasks, KG Off scores $\overline{\text{MPF}}_w=0.21$ (vs. equal-weight $\text{MPF}=0.34$), confirming that KG Off’s most damaging errors are also in the most sensitive properties. KG Smart retains $\text{MPF}_w=1.00$ (success-only). For standard tasks (Easy/Medium/Hard) where materials are well-known, MPF and MPF_w both approach 1.0 across all modes.

6.2 Results

table 4 presents aggregate results across the three conditions.

Table 4: Three-way KG ablation across 50 benchmark tasks with frozen knowledge graph ($n=50$ per mode). “Success-only” rows isolate output quality by excluding failed runs (which receive $\text{phys}=0.5$, $\text{MPF}=0$). 95% Wilson CIs for success rate.

Metric	KG Off	KG On	KG Smart
Success rate	100% [93,100]	94% [84,98]	100% [93,100]
<i>Overall (all tasks):</i>			
Physics score	0.853 ± 0.21	0.846 ± 0.24	0.933 ± 0.13
MPF	0.796 ± 0.30	0.752 ± 0.39	0.926 ± 0.16
<i>Success-only quality:</i>			
Physics score	0.853 ± 0.21	0.879 ± 0.19	0.933 ± 0.13
MPF	0.796 ± 0.30	0.801 ± 0.35	0.926 ± 0.16
Wall time (s)	12.9 ± 4.9	60.8 ± 73.2	28.7 ± 39.3
<i>By difficulty (success rate):</i>			
Easy (12)	100%	92%	100%
Medium (15)	100%	93%	100%
Hard (13)	100%	100%	100%
Novel (10)	100%*	90%	100%
<i>Novel tasks quality:</i>			
Physics score	0.590	$0.887^\dagger$	0.999
MPF	0.340	$0.889^\dagger$	1.000

*KG Off “succeeds” but fabricates properties ($\text{MPF}=0.34$).

† Success-only (9/10 tasks); N08 timed out.

Success rate. KG Off and KG Smart both achieve 100% success [93%, 100%] ($n=50$), while KG On reaches 94% [84%, 98%]. The difference between KG On and the other modes is not statistically significant ($p = 0.079$, Cohen’s $h = 0.50$), but the three KG On failures are systematic: N08 (Pyrathane, timeout), and two medium tasks where mandatory KG tool calls exhaust the iteration budget.

Output quality. KG Smart achieves the highest output quality: mean physics score of **0.933** and MPF of **0.926**, compared to 0.853 / 0.796 for KG Off (Cohen’s $d = 0.46$ for physics, a medium effect). KG On also outperforms KG Off on quality (0.879 / 0.801 success-only); the KG *does* help the agent select more accurate material properties.

Novel materials: the KG’s decisive edge. The clearest demonstration of KG value appears on the 10 novel-material tasks using fictional materials (Novidium, Cryonite, Pyrathane) whose properties exist only in the KG. KG Off “succeeds” on all 10 tasks but fabricates properties, producing a physics score of 0.590 and MPF of only 0.340. KG Smart achieves near-perfect quality (physics score 0.999, $\text{MPF} = 1.000$), while KG On scores 0.887 / 0.889 (success-only, 9/10 tasks). The knowledge graph is essential for materials absent from the LLM’s training data.

Wall time. KG Off is 2–5 $\times$ faster (12.9s mean) than KG Smart (28.7s) and KG On (60.8s), reflecting the iteration overhead of KG tool calls.




Figure 6: Success rate by difficulty level across three KG integration strategies. KG Smart (warm-start + lazy retrieval) matches KG Off on standard tasks while retaining high output fidelity. The “Novel” category (10 tasks) uses fictional materials whose properties exist only in the KG. Note: success rate alone is misleading for novel tasks; see the MPF/physics rows in table 4.

6.3 Analysis and Discussion

The three-way comparison reveals a clear trade-off: KG-free agents are fastest and most reliable, but KG-enabled agents produce higher-fidelity outputs. *KG Smart* balances both dimensions. Three mechanisms explain the mandatory KG pattern’s reliability penalty:

Iteration-budget exhaustion. KG On’s system prompt mandates `query_knowledge_graph` and `check_config_warnings` calls before every simulation attempt. These add 2–4 iterations, and for medium/hard tasks the agent can exhaust its iteration budget before

reaching `run_simulation`. KG Smart avoids this by deferring KG calls to the failure-recovery path.

Warning-induced conservatism. The `check_config_warnings` tool occasionally returns pre-emptive warnings (e.g. “CFL criterion may be violated”) that the model interprets as prohibitive, leading it to report a warning instead of proceeding. A warning that should trigger a *modification* instead triggers a *termination*.

Quality conditional on success. The key insight is that KG access *does* improve output quality when the agent completes the task. KG On’s success-only physics score (0.879) and MPF (0.801) both exceed KG Off (0.853 / 0.796). On novel materials, this gap widens dramatically: KG-enabled modes achieve $\text{MPF} \geq 0.89$ versus 0.34 for KG Off. The knowledge graph is not the bottleneck; the *integration pattern* is.

KG Smart: corrective retrieval + warm-start. Motivated by CRAG [25] (corrective RAG with adaptive retrieval) and AriGraph [26] (episodic KG memory for agents), we implemented a *KG Smart* integration combining:

1. **Warm-start injection.** Before the agent loop begins, we embed the task description with `nomic-embed-text` and query the Neo4j HNSW vector index for the top-3 most similar past successful runs. Their configurations are injected directly into the system prompt as few-shot reference examples, requiring no tool call.
2. **Lazy conditional retrieval.** KG tools remain available but the system prompt instructs the agent to use them *only* after a simulation failure or when material properties are genuinely unknown. This eliminates the mandatory KG-first workflow.

As table 4 shows, KG Smart achieves 100% success (vs. 94% for mandatory KG) while attaining the *highest* output quality across all modes (physics 0.933, MPF 0.926). It delivers the best of both worlds: KG Off reliability with KG-enhanced fidelity.

6.4 Physics-Aware Quality Metrics and Novel Material Validation

The preceding analysis showed that KG access improves output quality *conditional on success*, but standard success rate alone cannot capture this: a simulation with fabricated material properties still “succeeds” computationally. To expose this gap, we introduced three physics-aware metrics (section 6): Material Property Fidelity (MPF), Physics Score, and sensitivity-weighted MPF_w . These metrics measure how *correctly*, not just whether, the agent configures the simulation. We validate them on the 10 novel-material tasks, which provide the clearest test of KG utility: the three fictional materials’ properties exist *only* in the knowledge graph.

KG Smart achieves near-perfect scores ($\text{MPF} = 1.00$, physics = 0.999) across all completed novel tasks. The warm-start injection retrieves exact property values from the HNSW vector index before the agent loop begins, and the agent consistently uses them.

KG Off fabricates wrong properties for every material. table 5 shows the agent-chosen vs. ground-truth properties across all 10 novel tasks. KG Off assigns *different* wrong values each run (the LLM hallucinates non-deterministically), but the errors are consistently severe: for Pyrathane, k ranges from 0.15 to 0.35 W/mK, a $900\times$ to $2,080\times$ underestimate of the true $k=312$, effectively treating a refractory super-conductor as an insulator. In transient tasks, this causes output errors up to $|\Delta T_{\max}|=949$ K, a physically impossible result arising from numerical instability on a system made artificially stiff by the wrong conductivity.

Table 5: Material properties chosen by each mode across the 10 novel tasks in the 50-task ablation. KG Off fabricates different wrong values each run (ranges shown); KG Smart retrieves the ground-truth values exactly. The worst-case k error for Pyrathane is $2,080\times$.

Material	Mode	k (W/m K)	ρ (kg/m ³)	c_p (J/kg K)
Novidium	Ground truth	73.0	5420	612
	KG Off	10–50	3500–8960	130–900
	KG Smart	73.0	5420	612
Cryonite	Ground truth	0.42	1180	1940
	KG Off	0.15–35	960–8940	385–1370
	KG Smart	0.42	1180	1940
Pyrathane	Ground truth	312.0	3850	278
	KG Off	0.15–0.35	1600–1850	900–1470
	KG Smart	312.0	3850	278




Figure 7: Material Property Fidelity (MPF) per fictional material. KG Off scores below 0.5 for all three materials, with Pyrathane at 0.15 (conductivity error: $2,080\times$). KG Smart retrieves exact properties from the knowledge graph, achieving $\text{MPF} = 1.00$ across all materials.

Error propagation analysis. To quantify how wrong properties affect simulation outputs, we ran paired simulations: one with ground-truth properties (reference) and one with the fabricated properties from KG Off.

Table 6: Error propagation from fabricated properties (KG Off) to simulation output. Steady-state Dirichlet tasks (G1, C1, P1) show zero output error; transient/Robin tasks show significant deviations.

Task	Mat.	$ \Delta T_{\max} $ (K)	$ \Delta T_{\min} $ (K)	$ \Delta T_{\text{mean}} $ (K)	ϵ_k (%)	ϵ_ρ (%)	ϵ_{c_p} (%)
G1	Nov.	<0.1	<0.1	<0.1	86	45	18
G2	Nov.	<0.1	<0.1	20.7	38	65	78
G3	Nov.	<0.1	46.6	1.3	86	45	18
C1	Cryo.	<0.1	<0.1	<0.1	64	19	29
C2	Cryo.	<0.1	9.5	0.3	138	659	77
P1	Pyra.	<0.1	<0.1	<0.1	100	58	224
P2	Pyra.	949	<0.1	363	100	58	188

$\epsilon_p = |p_{\text{fab}} - p_{\text{true}}|/p_{\text{true}}$. P2’s k error (2,080 $\times$) produces 949 K $T_{\max}$ overshoot.

table 6 shows the results. Steady-state Dirichlet tasks (G1, C1, P1) are analytically independent of k , ρ , c_p , confirming zero output error regardless of property error, an important control. Transient and Robin tasks show significant deviations:

- **Pyrrhane P2:** $|\Delta T_{\max}|=949$ K ($T_{\max}=2,949$ K vs. 2,000 K reference); the 2,080 $\times$ k error causes numerical instability that *heats* the domain beyond its initial condition, a physically impossible result.
- **Novidium G3:** $|\Delta T_{\min}|=47$ K due to wrong conductivity affecting the Robin BC equilibrium.
- **Novidium G2:** $|\Delta T_{\text{mean}}|=21$ K from incorrect diffusivity altering the transient profile.




Figure 8: Error magnification: mean material property error (%) vs. output temperature error (%) for the four tasks with non-negligible output deviations. P2’s output error (24.4%) exceeds its input property error due to nonlinear amplification in the transient solver.

Sensitivity-weighted MPF. Using the sensitivity coefficients from the error propagation runs, MPF_w (defined in section 6) penalises errors in high-influence properties more heavily. Across the 10 novel tasks in the 50-task ablation, KG Off’s $\overline{\text{MPF}}_w=0.21$ (vs. equal-weight $\text{MPF}=0.34$), confirming that KG Off’s most damaging

Table 7: Cost-benefit of KG modes on novel tasks ($n=10$, 50-task ablation). Efficiency = $\text{MPF} / \text{wall-time}$ ($\times 10^{-2}$).

Mode	Succ.	Time (s)	Iter.	MPF	Phys. Score	Eff. ($\times 10^{-2}$)
KG Off	100%	11.7	3.0	0.34	0.59	2.92
KG On	90%	111.4	7.4	0.80	0.80	0.72
KG Smart	100%	58.0	4.9	1.00	1.00	1.72

KG Smart: +396% time, +194% MPF vs. KG Off; +2.4 $\times$ efficiency vs. KG On.

Algorithm 1 Adaptive KG Mode Selection

Require: Task description d , known-material list $\mathcal{M}$
Ensure: KG mode $m \in \{\text{Off, Smart-lazy, Smart-forced}\}$

```

1: if  $d$  contains explicit  $k$ ,  $\rho$ ,  $c_p$  values then
2:   return KG OFF
3: end if
4:  $\mu \leftarrow \text{EXTRACTMATERIAL}(d)$ 
5: if  $\mu = \emptyset$  then
6:   return KG OFF
7: else if  $\mu \in \mathcal{M}$  then
8:   return KG SMART (LAZY)
9: else
10:  return KG SMART (FORCED WARM-START)
11: end if

```

errors (particularly k for Pyrrhane and Cryonite) are concentrated in the properties with the highest sensitivity. KG Smart retains $\text{MPF}_w=1.00$ (success-only).

Cost-benefit analysis. table 7 compares the operational cost (time, iterations) against the accuracy benefit (MPF, physics score) for each KG mode. On novel tasks, KG Smart adds significant wall-time over KG Off (58.0s vs. 11.7s) but delivers $\approx 2.9\times$ higher MPF (1.00 vs. 0.34) and $2.4\times$ higher efficiency (MPF / wall-time) than KG On. KG On is slowest (111.4s on novel tasks) and its mandatory retrieval overhead yields diminishing returns: 90% success vs. 100% for both alternatives.

Adaptive KG decision framework. Based on the preceding analysis, we propose a simple decision algorithm (Algorithm 1) that selects the optimal KG mode for each task. Retrospective validation on all 50 benchmark tasks confirms high accuracy: it correctly assigns KG Off to tasks with explicit parameters, KG Smart (lazy) to tasks naming known materials, and KG Smart (forced warm-start) to novel-material tasks.

Cross-model validation. To verify that the KG’s value is model-agnostic, we repeated the novel-material experiment with `qwen3-coder-next` (80B total, 3B active MoE). table 8 shows that KG Smart achieves $\text{MPF}=1.00$ with *both* models, while KG Off fabricates wrong properties with both (different wrong values, but the same failure mode). This confirms that the warm-start vector search bypasses model-specific memorisation

Table 8: Cross-model validation on novel tasks ($n=7$, earlier experiment with v1 benchmark). KG Smart achieves perfect MPF with both LLM backends.

Model	Mode	Succ.	Qual.	MPF	Phys.
Qwen2.5-Coder 32B	KG On	57%	0.41	0.43	0.64
	KG Off	100%	0.61	0.34	0.63
	KG Smart	100%	0.96	1.00	1.00
Qwen3-Coder-Next 80B	KG On	71%	0.14	0.14	0.57
	KG Off	86%	0.54	0.40	0.70
	KG Smart	100%	0.96	1.00	1.00

Warm-start vector search retrieves exact properties before the LLM reasons, eliminating dependence on memorised material data.

gaps entirely.

Key takeaway. The KG’s value is not in making simulations *run* (LLMs can already do that) but in making simulations produce *correct physics*. For any domain with proprietary materials, novel compounds, or in-house measurement data, the KG Smart pattern transforms the system from an “expensive random number generator” into a reliable engineering tool. The error propagation analysis quantifies this: fabricated properties cause output errors up to 949 K (47% of the reference $T_{\max}$), while KG Smart eliminates these errors entirely by retrieving exact properties before the agent loop begins.

Narrowing the focus: correctness over completion. KG Off is the fastest mode and matches KG Smart on success rate for *standard* tasks with well-known materials. When the task description names “steel” or “copper”, the LLM’s parametric knowledge suffices and KG overhead is pure cost. However, for any deployment where output *correctness* matters, not just completion, KG-enabled modes are necessary. The preceding analysis makes this unambiguous: KG Off achieves MPF=0.34 and physics 0.59 on novel materials (fabricated properties), while KG Smart achieves MPF=1.00 and physics 0.999. The adaptive decision framework (Algorithm 1) encodes this distinction, routing tasks to KG Off when parameters are explicit and to KG Smart otherwise. We therefore focus the remaining analysis on **KG On vs. KG Smart**: both modes have access to the same knowledge base, yet KG Smart achieves 6 percentage points higher success and the highest quality scores. Understanding *why*, and which component (warm-start vs. lazy retrieval) contributes most, is the subject of the next section.

6.5 Failure Analysis: KG On vs. KG Smart

Having established that KG Smart matches KG Off at 100% success while outperforming both on quality, we trace KG On’s 3 remaining failures to their root causes. This analysis was performed *after* implementing an agent-level retry mechanism (see below) that eliminated stochastic failures; the 3 remaining failures are therefore all systematic.

Mitigating stochastic failures. In an earlier ablation run without retry logic, KG On exhibited 12/50 failures, of which 7 were stochastic (succeeded on re-run). We implemented an *auto-retry* mechanism directly in the agent: if the first attempt produces no `run_id` and used less than 60% of its iteration budget (indicating early stochastic exit rather than budget exhaustion), the agent automatically retries once. This reduced KG On failures from 12 to 3, eliminating all stochastic failures while not masking systematic ones.

KG On failure taxonomy. The 3 remaining KG On failures are all systematic:

Table 9: Failure mode classification for KG On’s 3 systematic failures (after auto-retry eliminates stochastic failures).

Failure mode	Count	Description
Budget exhaustion	2	Agent spends all iterations on repeated <code>query_knowledge_graph</code> and <code>check_config_warnings</code> calls without reaching <code>run_simulation</code> .
Timeout	1	N08 (Pyrathane): the mandatory KG query workflow combined with a numerically stiff problem exceeds the 420s time limit.

All 3 failures stem from the mandatory KG-first workflow: the agent must call `query_knowledge_graph` and `check_config_warnings` before *every* simulation attempt, consuming 2–4 iterations per cycle. For medium tasks, this exhausts the 25-iteration budget before `run_simulation` is called. For N08 (Pyrathane), the workflow inflates wall time beyond the timeout.

KG Smart: zero failures. KG Smart achieves 100% success with zero failures. The warm-start injection provides material properties and reference configurations *before* the agent loop, so the agent proceeds directly to `validate_config` → `run_simulation`. The average iteration count drops from 7.5 (KG On) to 4.2 (KG Smart), eliminating budget-exhaustion risk entirely. Even N08 completes successfully in 170s (5 iterations) because warm-start provides the correct Pyrathane properties upfront.

Decomposing warm-start vs. lazy retrieval. The failure data allows us to attribute KG Smart’s advantage to its two components without running a separate experiment. All 3 systematic failures stem from the mandatory *pre-simulation* KG workflow: the agent spends iterations on KG queries before ever attempting `run_simulation`. Warm-start injection eliminates this entirely: it provides material properties and reference configurations *before* the agent loop begins, so the agent can proceed directly to `validate_config` → `run_simulation`.

Lazy retrieval, by contrast, is exercised only on failure recovery: it allows the agent to query the KG *after* a simulation fails. Since budget-exhaustion tasks never

reach `run_simulation` at all, lazy retrieval cannot help them. We therefore conclude that **warm-start injection is the dominant factor** in KG Smart’s advantage over KG On: it directly prevents all 3 systematic failures (100% of KG On’s remaining failures after auto-retry). Warm-start also indirectly mitigates stochastic failures by reducing the iteration count (fewer decision points $\Rightarrow$ fewer opportunities for the LLM to produce a non-tool-call response). Lazy retrieval provides complementary value on simulation-error recovery but is secondary for the reliability gain.

7 Production Agent Quality Metrics

Beyond the controlled ablation, we analyse the system’s behaviour across all simulation runs accumulated during development, testing, and ablation experiments. Of these, 805 originate from a comprehensive automated parameter sweep across 8 physics studies (geometry, material, boundary conditions, initial conditions, time-stepping, 3D geometry, Robin convection, and multi-material) run via `scripts/sweep_full_study.py`; the remainder are interactive and ablation runs accumulated through normal system use.

Table 10: Agent ecosystem performance metrics from production database.

Metric	Value
Total simulation runs	1369
Overall success rate	97.8%
Unique agent tasks	110
First-try success rate	57.6%
Suggestion acceptance rate	0.0%
Avg. steps/task (simulation)	7.9
Avg. steps/task (database)	6.2
Avg. steps/task (analytics)	10.2
Avg. simulation time	1.5s

7.1 Success and Failure Modes

The 97.8% overall success rate (1,339/1,369 from PostgreSQL `simulation_runs`) demonstrates production-grade reliability for the full agent stack. Of the 2.2% failures, the majority are attributable to invalid boundary condition combinations (e.g. pure Neumann systems without a reference point), mesh resolution requests exceeding GPU memory, and numerical divergence on stiff transients.

7.2 Iteration Efficiency

Mean reasoning steps per completed task: Simulation Agent 7.9, Analytics Agent 10.2, Database Agent 6.2. The higher Analytics Agent count reflects its tendency to run multiple comparison queries before synthesising

a recommendation. The Database Agent’s lower count reflects the relative simplicity of structured SQL-style queries over natural-language reasoning.

7.3 First-Try Rate

57.6% of simulations run successfully on the first attempt (no debug/retry cycle). The remaining 42.4% require at least one modification-and-retry loop, typically for boundary condition corrections or CFL stability adjustments. This is consistent with the complexity of translating free-form natural language to precise numerical configurations.

7.4 KG Growth and Learning Curve

To isolate the KG’s causal effect on performance, we run a *controlled* growth experiment. The KG is initialised with only static material definitions (13 materials, 45 reference entries) but *zero* Run nodes, so the agent has no prior simulation experience. We then run 100 tasks (two shuffled passes through the 50-task benchmark) sequentially in KG Smart mode with writes enabled; each successful run adds one Run node with an HNSW-indexed embedding and up to five SIMILAR_TO edges. By pass 2, the KG contains 50 Run nodes from pass 1, providing warm-start context that was absent during pass 1.

To measure the effect, we perform a *paired comparison*: for each of the 50 benchmark tasks, we compare success-only MPF and physics score between pass 1 (zero prior runs) and pass 2 (50 prior runs), retaining only the 33 non-novel tasks that succeeded in both passes. Novel tasks are excluded because their fidelity depends on pre-seeded Material nodes rather than accumulated Run history. fig. 9 shows the result stratified by difficulty.




Figure 9: Paired KG growth comparison: success-only MPF and physics score for the same tasks run in pass 1 (0 prior Run nodes) vs. pass 2 (50 prior Run nodes). Only tasks succeeding in both passes are compared ($n=33$; novel tasks omitted as they rely on pre-seeded material definitions). Easy and medium tasks are at ceiling; **hard tasks gain +8.8% MPF and +5.8% physics score**, the category where similarity-based warm-start adds most value.

The experiment reveals a *stratified* growth effect. Easy

and medium tasks already achieve $\text{MPF} \geq 0.96$ in pass 1 because their parameters are either given explicitly or well-known to the LLM; accumulated run history cannot improve what is already near-perfect. Hard tasks, however, involve ambiguous descriptions, mixed boundary conditions, and tricky numerics where the agent benefits most from seeing similar prior runs: MPF rises from 0.783 to 0.870 (+8.8%), and physics score from 0.755 to 0.799 (+5.8%). The overall success rate is 88% (88/100), confirming that KG Smart remains robust even when starting without prior experience.

These results show that KG value is *difficulty-dependent*: static material knowledge covers simple cases, while accumulated run experience provides measurable gains on the hardest tasks where warm-start context disambiguates underspecified problems.

7.5 Comparison with Prior Work

FEM code generation. Bauer et al. [15] report $\approx 60\%$ one-shot success for GPT-4-driven FEniCS script generation; our 57.6% first-try rate is broadly comparable using a fully local, open-source model stack. Unlike their single-turn approach, PDE-Agents recovers from failures through multi-turn self-correction, achieving 97.8% final success. ALL-FEM [17] addresses the FEniCS automation problem through fine-tuning: a corpus of 1,000+ verified FEniCS scripts is used to fine-tune models from 3B to 120B parameters, and a multi-agent framework orchestrates problem formulation, code generation, and debugging. Their best model (GPT OSS 120B) achieves 71.8% code-level success on 39 benchmarks spanning elasticity, plasticity, fluid flow, and multiphysics, a substantially broader PDE scope than ours. PDE-Agents differs in three respects: (i) we use unmodified open-source LLMs without domain-specific fine-tuning, relying instead on tool-call orchestration and KG-augmented prompting; (ii) our system manages the full simulation lifecycle (configuration, execution, monitoring, debugging) rather than generating standalone scripts; and (iii) we provide a formal V&V study and a controlled ablation of knowledge graph integration modes, establishing when and why retrieval augmentation helps.

FEA benchmarks. FEABench [18] provides a systematic evaluation of LLM and agent capabilities for finite element analysis using COMSOL Multiphysics. Their best agent strategy generates executable API calls 88% of the time across problems in heat transfer, structural mechanics, and electromagnetics. Our work complements FEABench by targeting the open-source FEniCSx ecosystem and measuring not just execution success but *output quality* (physics score, material property fidelity), exposing cases where computationally “successful” simulations produce physically incorrect results due to fabricated material properties.

CFD agent systems. The computational fluid dynamics community has been particularly active in LLM-driven

automation. OpenFOAMGPT 2.0 [33] achieves 100% success across 450+ simulations using a four-agent pipeline (pre-processing, prompt generation, simulation, post-processing) with RAG-augmented retrieval. MetaOpenFOAM [34] decomposes CFD workflows into subtasks using MetaGPT’s assembly-line paradigm, reporting 85% pass rates at \$0.22 per case. These systems validate the multi-agent pattern we adopt but target a different solver ecosystem and do not evaluate knowledge graph integration or measure output fidelity against ground-truth material properties.

KG-augmented agents for materials science. Buehler [29] demonstrates retrieval-augmented ontological knowledge graph strategies for materials design with a fine-tuned MechGPT model, showing that graph-structured retrieval outperforms flat RAG by providing mechanistic relationships between concepts. Our KG Smart pattern shares the same insight (that structured knowledge improves over naïve retrieval) but applies it in a *live simulation loop* rather than a generative design setting, and our ablation study provides controlled evidence for when graph-augmented retrieval helps (novel materials) versus when it adds overhead without benefit (standard tasks with well-known properties).

8 System Limitations and Future Work

Solver scope. The current solver supports only scalar heat equations. Extension to vector mechanics (linear elasticity, Navier–Stokes) requires substantially richer tool schemas and system prompts.

Model scale and open-source evolution. The system defaults have been upgraded from the initial v1 models (qwen2.5-coder:32b, llama3.3:70b) to Qwen3-Coder-Next (80B total, 3B active MoE) [35] and Llama 4 Scout (109B total, 17B active MoE) [36], with cross-model validation confirming model-agnostic KG Smart performance (table 8). Qwen3-Coder-Next achieves higher config quality on standard KG Smart tasks (0.94 vs. 0.86) and perfect MPF on novel tasks, identical to the v1 model. Llama 4 Scout provides 10M-token context and native tool calling for the orchestrator and analytics agents. A systematic multi-model ablation across additional families (GLM-5, DeepSeek-R2) is planned as future work.

KG integration. The KG Smart pattern (section 6.3) achieves 100% success while producing the highest quality outputs, matching KG Off’s reliability. KG On retains a small gap (94%) due to budget exhaustion. The adaptive decision framework (Algorithm 1) provides practical deployment guidance, achieving 100% mode-selection accuracy. Future work will explore: (a) reinforcement learning from simulation outcomes to tune KG retrieval relevance, (b) dynamic confidence-based retrieval (only

query KG when the LLM’s uncertainty exceeds a threshold), and (c) richer warm-start context including failure diagnostics from similar past runs.

Verification scope. The V&V study covers only steady-state and simple transient heat equation benchmarks. A complete V&V programme would include higher-order elements, singularity-containing domains, and time-dependent convergence rates.

Parallelism and scalability. The current architecture runs agents sequentially. A task-level parallelism layer (e.g. running a parametric sweep as concurrent simulation agent calls) is under development.

Knowledge graph evolution. The 805-run parameter sweep has already created a dense graph with 4,000+ SIMILAR_TO edges. Future work will mine this corpus for IMPROVED_OVER relationships (did increasing mesh resolution reduce error without proportionally increasing wall time?), enabling the agent to autonomously suggest Pareto-optimal configurations. A second planned enhancement is citation traceability: agents citing specific ReferenceChunk IDs from indexed literature when justifying parameter choices, providing a fully auditable reasoning trail.

9 Conclusion

We have presented PDE-Agents, an open, containerised multi-agent ecosystem that automates finite element simulations through natural-language interaction with locally-deployed open-source LLMs. The system achieves 97.8% overall success across 1,369 production runs and demonstrates $\mathcal{O}(h^2)$ spatial convergence validated against analytical solutions.

The KG Smart integration pattern. Our three-way ablation study across 50 benchmark tasks with a frozen KG reveals that the *integration pattern*, not the knowledge source itself, determines KG utility. KG Off (no knowledge graph) achieves 100% task completion but fabricates material properties, producing MPF = 0.34 and MPF_w=0.21 on novel materials. KG On (mandatory KG access) retrieves correct properties but retains a 6% failure rate from budget exhaustion and timeout (even after implementing agent-level retry to eliminate stochastic failures). KG Smart, combining warm-start injection with lazy conditional retrieval, achieves **100% success with the highest output quality** (physics 0.933, MPF 0.926), resolving the reliability-vs-fidelity trade-off.

Failure analysis: why warm-start dominates. After implementing auto-retry to eliminate stochastic LLM failures, KG On’s 3 remaining failures (section 6.5) are all systematic: 2 budget-exhaustion cases and 1 timeout, all caused by the mandatory pre-simulation KG query

loop. Warm-start injection eliminates this bottleneck entirely by providing material properties and reference configurations before the agent loop begins, accounting for KG Smart’s 6-pp success rate advantage. Lazy retrieval provides complementary value on failure recovery but is secondary for the reliability gain.

KG value is difficulty- and domain-dependent. A controlled 100-task KG growth experiment shows that the KG’s benefit is concentrated where it matters most. Easy and medium tasks are already at quality ceiling (MPF $\geq$ 0.96) from the first pass; hard tasks, involving ambiguous descriptions and mixed boundary conditions, gain +8.8% MPF and +5.8% physics score from accumulated run history. For novel materials, the KG is indispensable: fabricated properties cause output errors up to 949 K, while KG Smart eliminates these errors entirely.

Lessons learned. Three design principles emerged from the ablation and failure analysis:

1. **Never make KG access mandatory in the agent’s critical path.** Mandatory pre-simulation retrieval creates fragile coupling between knowledge access and task completion. Front-loading context via embedding similarity (warm-start) decouples them.
2. **Front-load context, don’t force the agent to query for it.** The agent’s iteration budget is its scarcest resource. Injecting relevant context into the system prompt before the agent loop is strictly more efficient than requiring the agent to spend iterations formulating and executing KG queries.
3. **Let the agent decide when it needs more knowledge.** Lazy conditional retrieval respects the agent’s autonomy: it queries the KG only after a failure or when material properties are genuinely unknown, avoiding unnecessary round-trips on tasks where the LLM’s parametric knowledge suffices.

Broader implications. The warm-start + lazy-retrieval pattern is not specific to heat transfer simulation. Any tool-using agent domain with a growing knowledge base (materials science, drug discovery, circuit design, structural engineering) could benefit from the same architecture: embed the task, retrieve the most similar prior successes, inject them as few-shot context, and defer explicit knowledge queries to the failure-recovery path. As knowledge graphs grow and LLMs improve, the quality gap between KG-enabled and KG-free agents will only widen for novel or proprietary domains where the LLM’s training data provides no coverage.

The production metrics confirm that multi-turn self-correction is essential: the 57.6% first-try rate improves to 97.8% through iterative debugging, underscoring the value of agent-level orchestration over single-turn code generation. PDE-Agents establishes both a working platform and a rigorous empirical baseline for the emerging

field of LLM-driven autonomous simulation, demonstrating that structured knowledge integration transforms an agent from an “expensive random number generator” into a reliable engineering tool.

Data Availability

All code, evaluation scripts, and result JSON files are available at <https://github.com/MatPro-IFE/pde-agents>. The Neo4j knowledge graph seed data is included in the repository at `knowledge_graph/seeder.py`.

Author Contributions

Sayan Adhikari: Conceptualization, Methodology, Investigation, Software, Data Curation, Formal Analysis, Visualization, Writing – Original Draft, Writing – Review & Editing. **Gulshan Sharma:** Conceptualization, Methodology, Investigation, Validation, Writing – Review & Editing. **Øyvind Jensen:** Supervision, Writing – Review & Editing.

Declaration of Generative AI in the Writing Process

During the preparation of this manuscript, Anthropic’s Claude Opus 4.6 was used to assist with language editing, grammar refinement, and the programmatic generation of TikZ figures (Figures 1 and 2). All scientific content, experimental design, analysis, and interpretation were performed by the human authors, who reviewed and take full responsibility for the final manuscript.

Competing Interests

The authors declare no competing interests.

Acknowledgements

The authors thank the anonymous reviewers for their constructive feedback. Computational resources were provided by the Institute for Energy Technology (IFE).

References

- [1] Tom B. Brown, Benjamin Mann, Nick Ryder, et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901, 2020.
- [2] Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35, 2022.
- [3] Shunyu Yao, Jeffrey Zhao, Dian Yu, et al. ReAct: Synergizing reasoning and acting in language models. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2023.
- [4] Isaac E. Lagaris, Aristidis Likas, and Dimitrios I. Fotiadis. Artificial neural networks for solving ordinary and partial differential equations. *IEEE Transactions on Neural Networks*, 9(5):987–1000, 1998.
- [5] Maziar Raissi, Paris Perdikaris, and George E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019. doi: 10.1016/j.jcp.2018.10.045.
- [6] George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, et al. Physics-informed machine learning. *Nature Reviews Physics*, 3(6):422–440, 2021.
- [7] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, et al. Fourier neural operator for parametric partial differential equations. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [8] John Jumper, Richard Evans, Alexander Pritzel, et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873):583–589, 2021. doi: 10.1038/s41586-021-03819-2.
- [9] Patrick Lewis, Ethan Perez, Aleksandra Piktus, et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474, 2020.
- [10] Darren Edge, Ha Trinh, Newman Cheng, et al. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024.
- [11] Andres M. Bran, Sam Cox, Oliver Schilter, et al. ChemCrow: Augmenting large-language models with chemistry tools. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- [12] Yubo Ma, Zhibin Liu, Liangming Pan Liang, et al. SciAgent: Tool-augmented language models for scientific reasoning. *arXiv preprint arXiv:2402.11451*, 2024.
- [13] Anders Logg, Kent-Andre Mardal, Garth N. Wells, et al. *Automated Solution of Differential Equations by the Finite Element Method: The FEniCS Book*. Springer, 2012. doi: 10.1007/978-3-642-23099-8.
- [14] Igor A. Barrata, Joseph P. Dean, Jørgen S. Dokken, et al. DOLFINx: The next generation FEniCS problem solving environment. *Zenodo*, 2023. doi: 10.5281/zenodo.10447666.

- [15] Philipp Bauer, Patrick Henning, and Janna Schaeffers. Large language models as automatic generators of FEniCS code for solving partial differential equations. *arXiv preprint arXiv:2312.09801*, 2023.
- [16] Wei Jiang, Keyi Chen, Minghan Wang, et al. LLM4FEM: Leveraging large language models for finite element method. *arXiv preprint arXiv:2405.03719*, 2024.
- [17] Rushikesh Deotale, Adithya Srinivasan, Yuan Tian, Tianyi Zhang, Pavlos Vlachos, and Hector Gomez. ALL-FEM: Agentic large language models fine-tuned for finite element methods. *arXiv preprint arXiv:2603.21011*, 2026.
- [18] Nayantara Mudur, Hao Cui, Subhashini Venugopalan, Paul Raccuglia, Michael P. Brenner, and Peter Norgaard. FEABench: Evaluating language models on multiphysics reasoning ability. *arXiv preprint arXiv:2504.06260*, 2025.
- [19] LangChain AI. LangGraph: Build stateful, multi-actor applications with LLMs, 2024. URL <https://github.com/langchain-ai/langgraph>.
- [20] Qingyun Wu, Gagan Bansal, Jieyu Zhang, et al. AutoGen: Enabling next-generation LLM applications via multi-agent conversation. In *Proceedings of EMNLP Industry Track*, 2023.
- [21] João Moura. CrewAI: Framework for orchestrating role-playing, autonomous AI agents, 2024. URL <https://github.com/joaoandmoura/crewai>.
- [22] Xinyi Liao, Hao Zhang, and Yutao Chen. Retrieval-augmented generation for engineering design documentation. *arXiv preprint arXiv:2307.04512*, 2023.
- [23] Yujia Gao, Shang Liu, Peng Shi, and Jimmy Lin. Retrieval-augmented code generation for universal information extraction. *arXiv preprint arXiv:2311.02555*, 2023.
- [24] Zheng Yang, Wenyan Li, and Peng Zhang. Simulation parameter suggestion via retrieval-augmented generation. *arXiv preprint arXiv:2403.09512*, 2024.
- [25] Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective retrieval augmented generation. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024. arXiv:2401.15884.
- [26] Petr Anokhin, Nikita Kornaev, Andrey Babkin, and Aleksandr I. Panov. AriGraph: Learning knowledge graph world models with episodic memory for LLM agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. arXiv:2407.04363.
- [27] Vineeth Venugopal, Soumya Sahoo, Gurinder Agastya, et al. MatKG: The largest knowledge graph in applied materials science. *arXiv preprint arXiv:2209.11632*, 2022.
- [28] Casper W. Andersen, Rickard Armiento, Evgeny Blokhin, et al. OPTIMADE: Towards an open database for computational materials science. *Scientific Data*, 8(1):217, 2021. doi: 10.1038/s41597-021-00974-z.
- [29] Markus J. Buehler. Generative retrieval-augmented ontologic graph and multiagent strategies for interpretive large language model-based materials design. *ACS Engineering Au*, 4(2):241–277, 2024. doi: 10.1021/acsengineeringau.3c00058.
- [30] Christophe Geuzaine and Jean-François Remacle. Gmsh: A 3-d finite element mesh generator with built-in pre- and post-processing facilities. *International Journal for Numerical Methods in Engineering*, 79(11):1309–1331, 2009. doi: 10.1002/nme.2579.
- [31] IBM Research. Docling: Document processing for AI, 2024. URL <https://github.com/DS4SD/docling>.
- [32] ASME. Guide for verification and validation in computational solid mechanics. Technical Report ASME V&V 10-2006, American Society of Mechanical Engineers, 2006.
- [33] Hernan Chen, Luca Mangani, and Gabriel Casas. OpenFOAMGPT 2.0: End-to-end, trustworthy automation for computational fluid dynamics. *arXiv preprint arXiv:2504.19338*, 2025.
- [34] Yuxuan Chen, Xu Zuo, Yifei Yang, et al. MetaOpenFOAM: An LLM-based multi-agent framework for CFD. *arXiv preprint arXiv:2407.21320*, 2024.
- [35] Ruisheng Cao, Mouxiang Chen, Jiawei Chen, Zeyu Cui, Yunlong Feng, Binyuan Hui, Yuheng Jing, Kaixin Li, Mingze Li, Junyang Lin, Zeyao Ma, Kashun Shum, Xuwu Wang, Jinxi Wei, Jiaxi Yang, Jiajun Zhang, Lei Zhang, Zongmeng Zhang, Wenting Zhao, and Fan Zhou. Qwen3-coder-next technical report. *arXiv preprint arXiv:2603.00729*, 2026.
- [36] Meta AI. The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation. <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>, 2025. Accessed 2026-04-15.